

Causal Inference in Population Trends: Searching for Demographic Anomalies in Big Data

March 30, 2021

Abstract

The proliferation of big data, wider access to advanced computing platforms, and the development of powerful statistical algorithms can uncover hidden anomalies in social data, previously dismissed as noise. Here, we demonstrate how researchers can combine causal inference techniques and abductive reasoning using an example that identifies fertility and mortality anomalies on twenty years of complete demographic data in the United States. We uncover real, “hidden” baby booms/busts and mortality spikes/dips, distinguishable from regular trend variations. We identify more than 22 fertility and 156 mortality anomalies, totaling more than 200k anomalous births and 600k anomalous deaths. Notable detectable mortality anomalies include the September 11 2001 terrorist attack in New York and the emergence and acceleration of the opioid epidemic in New Hampshire. Notable fertility anomalies include the “missing births” in Louisiana after Hurricane Katrina and the reduction in fertility behavior after the September 2008 stock market crash in Connecticut, among others. The combined causal inference and abductive reasoning approach can be readily adapted to find other, undiscovered social phenomena or to evaluate the efficacy of important public policies.

Keywords: mortality, fertility, causal inference, abductive reasoning, big data

1 Introduction

The search for causality is one of the primary motivations of social scientists (??). Causal inference, the identification of causal effects, attempts to uncover the underlying mechanisms that result in changes in a phenomenon and has a long history as an approach to inquiry in the social sciences (?). Often, the identification of a casual mechanism is accomplished via a randomized control trial (RCT) or expert knowledge applied to a natural experiment (?).

Unfortunately, both RCTs and natural experiments have drawbacks. RCTs are often prohibitively expensive for many scientists, precluding their widespread adoption in the social sciences (?). RCTs are also sometimes impossible to deploy in the “real world.” For instance, an RCT to study country-level austerity measures would be difficult, if not impossible, to implement. RCTs to contraceptive practices would be immoral. Natural experiments are a cheaper alternative but provide less certainty about causation, fewer opportunities to replicate work, and require expert knowledge of the natural experiment to be properly employed (?).

Additionally, both RCTs and natural experiments are based on traditional hypothesis-based testing, the deductive scientific paradigm where a question is first posed and scientists generate or find data to answer the question. Such approaches, while extensively utilized and time-tested, increase the likelihood that a researcher will miss important phenomena that might go unnoticed because it falls outside the hypothesis, or the study could suffer from extensive p-hacking (???). At the same time, purely inductive data mining can result in false assertions or spurious findings, especially in the hands of those who are not field experts (?). Abductive reasoning, sometimes described as “inferring cause from effect” (?), moves the scientific process from the inductive to the deductive, and sometimes back and forth, to reach conclusions (?). Physical scientists use abductive reasoning more commonly than social scientists. The discovery of Neptune in 1846 based on unexplained perturbations of the orbit of Uranus (?), the discovery of galactic migration of planets (?), and the current generation of gravitational wave detectors such as the Laser Interferometer Gravitational-Wave Observatory (LIGO; ?) demonstrate the power of the abductive approach in the physical sciences.

In a data-rich environment, abductive reasoning can be more easily applied to first find effects and then infer their cause ex-post, with careful, measured, additional research. The modern proliferation of data, advances in high performance (“super”) computing, and the development of powerful statistical algorithms provide an environment ripe for wedding causal inference with abductive reasoning (??), with the potential for researchers to find hidden or understudied social phenomena (?). While the availability of big data and high performance computing allows novel exploration of data through causal inference (???), relatively few studies utilize causal inference techniques in the study of demographic phenomena. However, understanding social phenomena using these advances reveals important insights into society (??) and allows us to better monitor population trends (??). The rich demographic data available in the United States make the potential revelation of interesting and important demographic phenomena not only possible, but extremely plausible – even if identification of the phenomena initially occurs without identifying the underlying cause.

In this paper, we combine a modern statistical outlier detection algorithm (?), abductive reasoning, and large demographic data sets on more than 20 years of mortality and fertility trends to identify anomalous demographic behavior and demonstrate the power of combining these approaches. In essence, we identify *effects without knowledge of the cause*. We ask one fundamental question regarding demographic anomalies: What are the hidden baby booms/busts and mortality spikes/dips in the United States over the last twenty years? We do not necessarily know the causes of these anomalies and suggest causes for select anomalies. Even without knowledge of the causes, simple identification of anomalies allows scientists with more detailed knowledge of local population dynamics, state-level policy making, or macro-economics a starting place to begin further investigation. From explanations of phenomena that may have previously gone unnoticed, social scientists may be able to better forecast populations and provide policy solutions for impending problems such as climate change.

2 Abductive Reasoning and Causal Inference

Scientific inquiry falls under one of three modes of logical reasoning: deduction, induction, and abduction. Deductive reasoning starts with a theory and makes a prediction of what the observations should be if that theory is correct. The classic example of deductive reasoning essentially goes (a) “all men are mortal,” (b) “Socrates is a man,” (c) “therefore, Socrates is mortal.” Deductive reasoning has a long history in the sciences and most hypothesis testing in the social sciences follows this logical reasoning. For example, John Snow used deductive reasoning to stop the 1854 London cholera outbreak. He hypothesized that cholera was waterborne and deduced that a bad water pump was responsible for the outbreak.

Inductive reasoning is directionally opposite of deductive reasoning. Rather than starting with a theory and making a prediction, induction starts with an observation and makes a theory based on the likeliest reason. For example, the Hispanic Mortality Paradox first appeared in the scientific literature in the mid-1980s (?) as an observation that Hispanics in the United States had better health outcomes than expected given their socioeconomic status. Based on this observation, numerous theories arose to explain this paradox.

Both deductive and inductive reasoning are well established in science but are not without their shortcomings. In both approaches, it is impossible, or at least improbable, for the premise to be true and the conclusion false. Furthermore, if a hypothesis is rejected under a deductive or inductive approach, there is no alternative explanation – there is either the presence or absence of support.

Abductive reasoning is a third, less well-known way that reasons the best explanation given the observations (?). It is most concerned with generating and evaluating explanatory hypotheses (?) by looking for a pattern and suggesting a hypothesis (??), then testing the hypothesis. Unlike the more rigid deductive and inductive reasoning, abductive reasoning allows the scientist to move between deduction and induction and fits well during exploratory data analysis (???). During exploratory data analysis, abductive reasoning helps find a pattern and suggest a hypothesis, deduction helps refine and build a testable hypothesis, and induction is the empirical support (?).

Abductive reasoning and exploratory data analysis are best employed in a data rich en-

vironment – the very data environment social scientists and demographers find themselves in the present moment. In a data rich environment, a hypothesis is a hidden liability (?) because it focuses our attention on only that which we seek. Yanai and Lercher (?) ask if focusing on a specific hypothesis prevents new discovery. By way of illustration, they created a simple dataset that if plotted would reveal an image of a gorilla waving. Students instructed to find support for specific hypotheses were less likely to discover the gorilla compared to those who were simply asked “What do you conclude from the dataset?”. Focusing on specific hypotheses often prevents us from finding other interesting phenomena. After all, it was the publication of the Hispanic paradox itself that has generated more than thirty years of devoted scholarship. At the same time, when paired with causal inference techniques, abductive reasoning provides an empirical platform for the beginning point of discovery of potentially interesting social phenomena; it provides an important tool for sifting large amounts of data to find interesting results without the blinders of hypothesis testing and potentially helps us to ask questions that we would not previously think to ask.

3 Materials and Methods

3.1 Method

We use a statistical time series outlier detection algorithm (?), implemented in the R programming language (?) via the tsoutliers package (?). This algorithm iteratively uses ARIMA models to 1) identify potential outliers or anomalies and 2) refit the ARIMA with the outliers removed to produce a counter-factual time series. Here we briefly summarize and describe the method.

Often, the behavior of a time series can be described and summarized in ARIMA models. If a series of values, y_t^* , is subject to m interventions or outliers at time points t_1, t_2, \dots, t_m with weights ω then y_t^* can be defined as

$$y_t^* = \sum_{j=1}^m \omega_j L_j(B) I_t(t_j) + \frac{\theta(B)}{\phi(B)\alpha(B)} a_t \quad (1)$$

Where $I_t(t_j)$ is an indicator variable with a value of 1 at observation t_j and where the j th outlier arises,

$\phi(B)$ is an autoregressive polynomial with all roots outside the unit circle,
 $\theta(B)$ is a moving average polynomial with all roots outside the unit circle,
and $\alpha(B)$ is an autoregressive polynomial with all roots on the unit circle.

We examine three types of outliers at time point t_m :

1. additive outliers (AO), defined as $L_j(B) = 1$;
2. level shift outliers (LS), defined as $L_j(B) = 1/(1 - B)$; and
3. temporary change outliers (TC), defined as $L_j(B) = 1/(1 - \delta B)$ where δ is equal to 0.7.

Colloquially, additive outliers arise when a single event causes the time series to unexpectedly increase/decrease for a single time period; level shift outliers arise when an event causes the time series to unexpectedly increase/decrease for multiple time periods; and temporary change outliers arise when an event causes the time series to unexpectedly increase/decrease with lingering effects that decay over multiple time periods.

An outlier is detected with the estimated residuals using a regression equation

$$\pi(B)y_t^* \equiv \hat{e} = \sum_{j=1}^m \omega_j \pi(B) L_j(B) I_t(t_j) + a_t \quad (2)$$

where $\pi(B) = \sum_{i=0}^{inf} \pi_i B^i$.

Equations ?? and ?? allow for an automatic detection of anomalies iterated over a two-stage process.

In stage 1, anomalies are located. First, an ARIMA model is fit to the time series using the `forecast` package in R (??) where the best performing ARIMA model is selected based on the Bayesian information criterion (BIC). Next, the residuals from the forecast are checked for their significance using equation ?? where only anomalies above a critical t -static are considered “true” anomalies ($|\tau| \geq 3.5$). We choose this critical value based on Chen and Liu’s (?) recommendation to minimize Type I or false-positive anomalies. Finally, two additional rules are implemented: If multiple anomalies are detected at the same time point, only the most significant anomaly is selected and if anomalies of the same type at consecutive time periods are detected, only the anomaly with highest t -statistic is selected.

In stage 2, anomalies are removed from the time series and a new ARIMA model is chosen and fit. The selection of the initial ARIMA model could have been affected by the

presence of the anomalies, making some anomalies spuriously identified. To correct for this, a new ARIMA model is fit accounting for additional regression effects in equation ?? from the list of candidate anomalies identified in stage 1, effectively removing the anomalies from the time series. Each anomaly is then reassessed under the new model and those anomalies that are no longer significant are removed.

These two stages are then iterated until no additional anomalies are detected.

3.2 Toy example

?? shows a toy example for anomaly detection in a time series using the example of Hurricane Katrina in Orleans Parish. On August 23 2005, Hurricane Katrina, a category 5 hurricane, struck southern Louisiana causing widespread damage and destruction in Orleans Parish in particular. The displacement from the Hurricane and the federal response were well documented (??) and Census estimates suggest Orleans Parish lost more than 200,000 residents between 2005 and 2006. What would have been Orleans’ population estimate had Hurricane Katrina *not* occurred? A simple counter-factual estimate might keep population size just under 500,000 people (\hat{y}).

In the Hurricane Katrina example, we have knowledge of the impact of Hurricane Katrina after the fact or ex-post-facto in order to create the counter-factual time series \hat{y} . But is this anomaly detectable without knowledge of the Hurricane Katrina? In other words, can we detect the reduction of Orleans Parish’s population between 2005 and 2006 using only the time series? Absolutely. The tsoutliers package identifies 2006 as an extremely strong level-shift outlier (t -stat = -96.39). The real world is hardly this simplified where a direct intervention is known, is large, and is testable.

3.3 Data

We search for demographic anomalies using the Center for Disease Control and Prevention’s online WONDER monthly fertility (2003-2017) and mortality databases (1999-2016) for all fifty states and the District of Columbia (??). These data sets contain every birth and death record in the United States over the time periods of interest, representing the universe of both mortality and fertility data in the US. These data are considered the “gold standard”

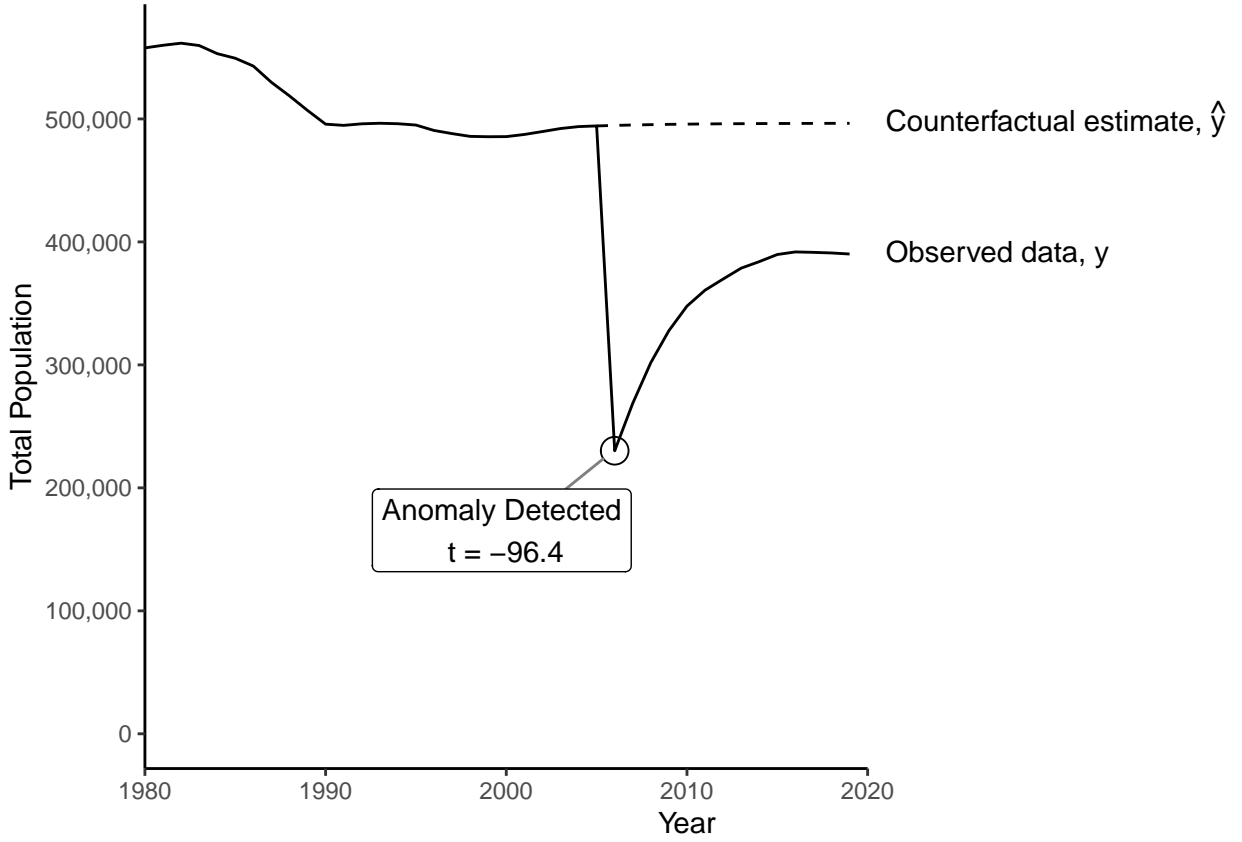


Figure 1: **An example of anomaly detection using Hurricanes Katrina and Rita in Orleans Parish Louisiana in 2005.** Hurricane Katrina struck Louisiana in 2005 and we use it as a toy example to illustrate our approach. This figure shows the annual time series of total population in Orleans Parish between 1980 and 2019. Between 1990 and 2005, Orleans Parish total population changed from 495k to 494k, suggesting a possible ‘plateau’ in the population (illustrated with the dotted ‘counterfactual’). Hurricane Katrina and the widespread population loss of more than 200k people represent a very strong anomaly ($t=-96.39$).

of data collections (?) and have been considered “complete” since 1968 (?). We search over each state equivalent’s (n=51) mortality (n=228) and fertility (n=180) monthly time series for a total of 20,808 state-months of data.

3.4 Robustness Check

Chen and Liu (?) report results of a simulation exercise of varying critical value levels. They simulated various anomaly-free time series and injected specific anomalies of varying σ levels. For a critical value of 3.5 – the level Chen and Liu recommend to minimize both false-positive (Type I) and false-negative (Type II) error rates – they report Type I error rates between 0.0% and 0.1% and Type II error rates or the statistical power between 40% and 60% with injected anomalies equal to 4σ . We follow their lead by simulating time series to test for Type I and Type II errors. We base our simulation on the 41 time series that do not contain any anomalies (see **Results** below; 33 without a single fertility anomaly and 8 without a single mortality anomaly).

To test for false-positive rates, we simulate 1000 time series for each of the 41 state-components (fertility or mortality) using the fitted ARIMA model for each state-component. This allows the simulated time series to mimic the underlying true time series (see **Supplementary Material** for the underlying code). Any anomalies we detect in the 7,764,000 candidate time steps are spurious anomalies. We detect 6,588 anomalies for a false-positive error rate of 0.085%.

To test for false-negative rates, we use the underlying true, sans-anomaly time series but introduce a single anomaly (magnitude equal to 3.5σ) at the same time point in each series. If the `tsoutliers` algorithm fails to detect the anomaly, it would be considered a false-negative error. We detect 27,904 out of 41,000 candidate time series giving us a false-negative error rate of 31.9%. This provides a statistical power of 68.1%.

Based on these tests we think it is clear that our probability of committing a Type I error (or claiming an anomaly is true when in fact is not) is very low. But, based on our false-negative error rate, it is possible that the anomalies we do detect are *conservative*. Our statistical power is of just 68.1% suggests a higher than normal likelihood that true effects are going undetected. We are unlikely to commit an error of commission (ie, including

a false anomaly) but are more likely to commit an error of omission (ie omitting true anomalies). Thus our results should be considered conservative.

3.5 Replication Files

All data and code necessary to reproduce the reported results are licensed under the CC-BY-4.0 license and are publicly available as https://osf.io/yc3r4/?view_only=3ce3ca267c7047e0942940ebd08846d0. The analyses were performed in *R* (?).

4 Results

We detect numerous anomalous mortality and fertility events at the US state-level since 1999. A full listing of these anomalies can be found in the **Supplementary Materials**. We begin by highlighting significant fertility and mortality anomalies across all three types of outliers (Additive Outliers, Level Shift Outliers, and Temporary Change outliers) with plausible explanations. We then highlight two strong anomalies that belie explanation. Finally, we conclude with a summary of the anomalies we detect.

4.1 Fertility

We begin with a clear temporary change (TC) outlier for fertility in Louisiana (??a). Here we detect two TC outliers, nearly back-to-back in 2005 (2005:08 $t = -6.002$; and 2005:10 $t = -4.997$), coinciding with the destruction of Hurricane Katrina from the same year. The migration associated with Hurricane Katrina has garnered most of the attention of social scientists (??), but the hurricane had a very clear impact on fertility behaviors too, creating a mini “baby bust” in Louisiana with the departure of many people in their childbearing years. We estimate 4,411 fewer births in Louisiana compared to the counter-factual, likely attributable to the Hurricane.

By way of illustration, we also highlight a level shift (LS) outlier for fertility in Connecticut. Here we detect a strong ($t = -3.71$ for -285 births/month) level shift toward lower fertility beginning in August 2009 (2009:08) that continues through the end of the period (??b). This LS toward lower fertility reduced the number of births by 28,752 or 9% lower

than the counter-factual time series. What is the cause? One possibility is the stock market crash of 2008, where the Dow Jones Industrial Average had the largest single-day loss up to that point, occurred just 11 months before we detect a shift toward lower fertility. It is possible that the trend toward lower fertility in August 2009 and the stock market crash in September 2008 are linked. TO determine this with any certainty, more research is needed. We are speculating only to illustrate how causal inference could work in an abductive research agenda. Once an anomaly is found, its discovery serves as a spur to further research. A potentially interesting future line of research would be to link fertility anomalies to the kinds of economic shocks that may lead to out-migration of those in their childbearing years.

4.2 Mortality

For mortality, we will use New York State as an example (??a). We identify seven anomalies in the mortality time series for New York, all with t -statistics in excess of 3.91, making these anomalies unlikely to be due to chance. The algorithm correctly identifies September 2001 as an additive outlier (2001:09 $t= 6.396$) where there were 1,628 more deaths in that month than anticipated. This mortality event is likely caused by the September 11 terrorist attack on the World Trade Center that immediately killed 2,606 people and the detection of this mortality event provides confidence in our detection of other anomalies.

In ??a, notice the strong level shift (LS) that occurs in February 2004 (2004:02 $t=-5.594$) which prevented 869 deaths per month. This shift totals more than 144,000 averted deaths compared to the counter-factual time series and is the single largest mortality protective anomaly among all states. This translates to 7% fewer deaths than expected over the time period. What is driving this mortality protection? What policies did NY put into place that might have contributed to this considerable mortality reduction? What environmental or economic conditions may have changed? These are the kinds of questions that arise from our analyses. The purpose of our paper is not to answer these questions, but our findings underscore a way of initially approaching social science research that takes advantage of big data techniques and new data availability that may lead us to answer new questions. By identifying anomalies through an inductive process, researchers can then

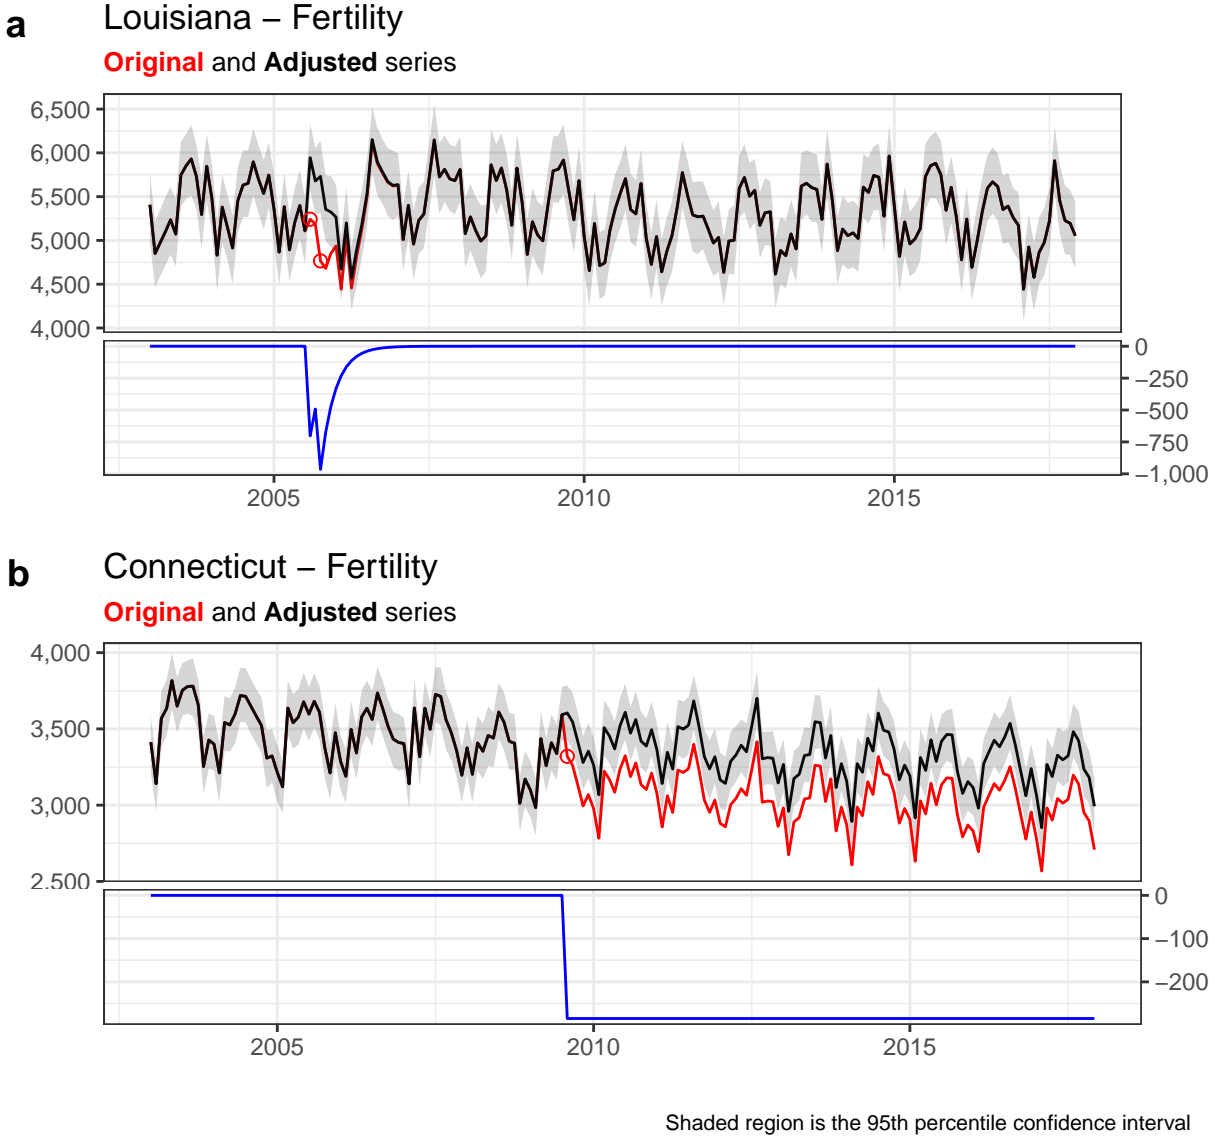


Figure 2: **Anomaly detection in Louisiana (a) and Connecticut (b) state fertility, 2003-2018.** The top part of each panel contains the original time series (light gray), the adjusted, counter-factual time series in the absence of anomalies (blue), and the red dots correspond to the onset of detected anomalies. The bottom part of each panel contains the magnitude and type of the outlier in red. In (a), we detect two outliers, back-to-back, likely resulting from Hurricane Katrina in August and October 2005, representing a decrease of more than 4,400 births due to the hurricane (t -statistics= -6.002 and -4.997). In (b), we detect one outlier, a level shift outlier (LS) in August 2009. We believe this reduction is attributable to the stock market crash 11 months earlier (t -statistic = -3.71).

look for underlying causes. Once those causes are identified, researchers can then use a deductive process to see if such events predict other (or future) anomalies.

To see the potential for combining abductive reasoning with causal inference, contrast the mortality protection in New York with the enhanced mortality in New Hampshire (??b). In New Hampshire we detect two significant level shifts (LS) in the monthly mortality data, first in April 2010 and again in November 2014 (2010:04 $t=3.51$; 2014:11 $t=6.68$). These anomalies suggest New Hampshire experienced 8,159 more deaths (+9% more than expected) in a seven-year period beginning in early 2010. These are events not experienced by neighboring states during the same time period, and this is the single largest percentage mortality increase/decrease we detected among all states. Not coincidentally, NH has the second highest opioid-related mortality in the US (?) and it is likely that we detect this epidemic in our results. Isolating these anomalies and testing them against opioid sales data might yield intriguing results.

4.3 Interesting Anomalous Fertility/Mortality events

In the examples above, we highlighted four fertility/mortality anomalies with plausible explanations. In the case of New York and New Hampshire, the mortality anomalies have plausible explanations. It seems likely that New York’s AO anomaly in September 2001 is caused by the 9/11 tragedy and the rise in New Hampshire’s mortality starting in 2010 could be linked to the opioid epidemic. Similarly, Louisiana’s TC anomalies seem linked to Hurricanes Katrina and Rita while the LS anomaly in Connecticut’s fertility appears linked to the Great Recession. However, we detect numerous other demographic anomalies in other states, on the causes of which we will not speculate. ?? shows two such unexplained anomalies.

In ??a we identify a single additive anomaly in Hawaiian fertility in May 2014. This is a strong anomaly with a t -statistic of 4.96, 15% above the counter-factual time series. This single, anomalous month is also the second highest monthly births in the time series. We have no plausible explanation for this anomaly. We do not believe this is simply a data error as the other extreme values, September 2008 and February 2005 with the highest and lowest recorded fertility respectively, were not identified as anomalous events. Even if we

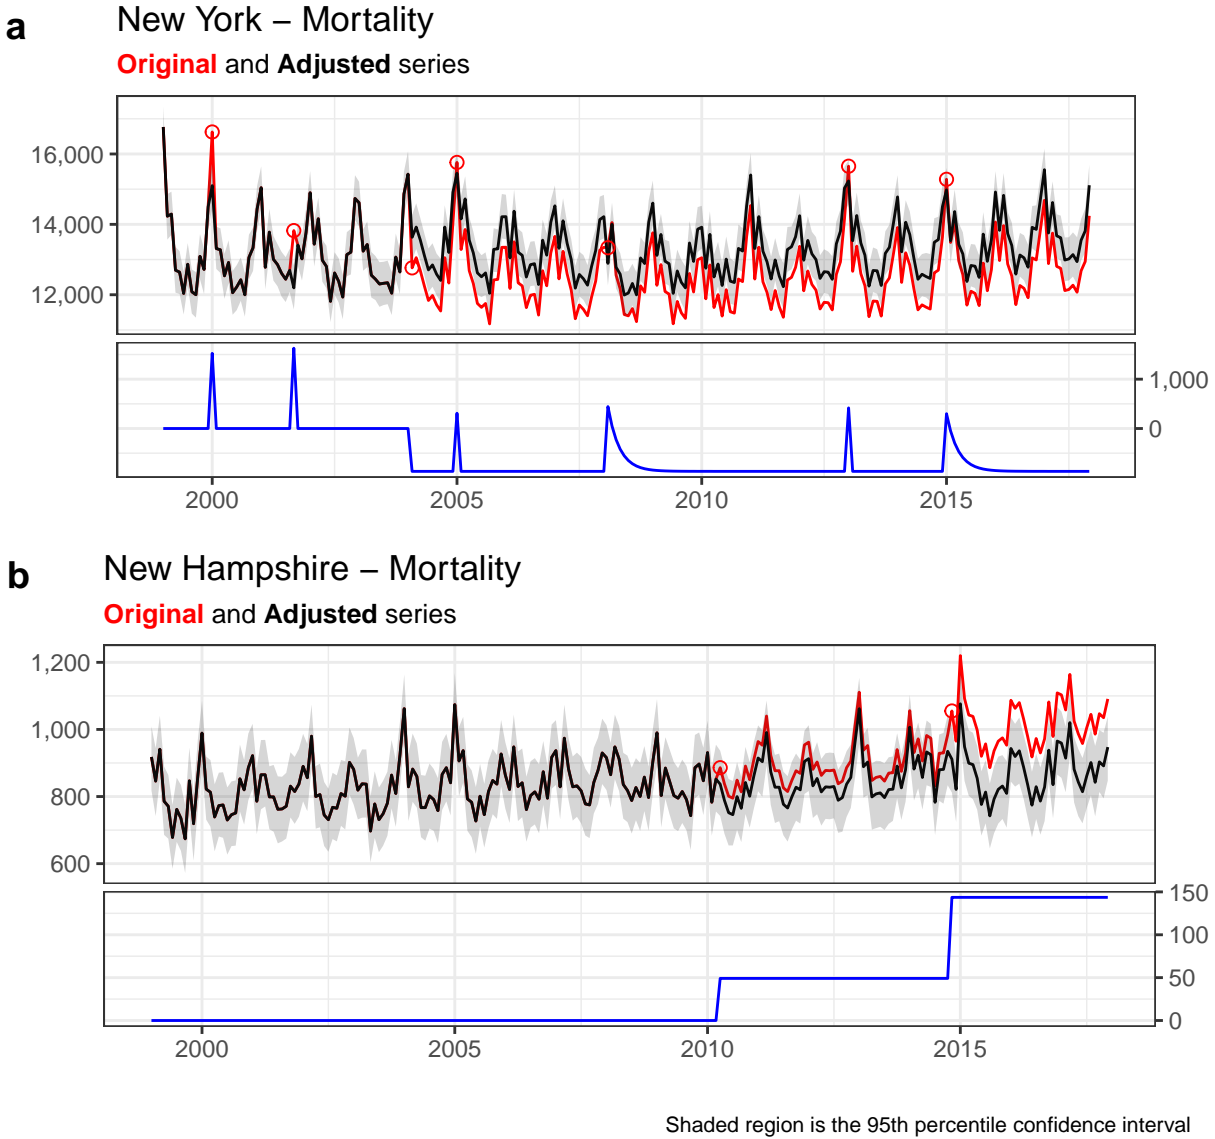


Figure 3: **Anomaly Detection for New York (a) and New Hampshire (b) state mortality, 1999-2016.** In New York (a), we detect additive outliers (AO) in January 2000, September 2001, January 2005, and January 2013; temporary change (TC) outliers in February 2007 and January 2015; and a level shift (LS) starting in February 2004. In New Hampshire (b), we detect two outliers, both level shift outliers (LS) in April 2010 and again in November 2014. These anomalies suggest New York experienced a significant mortality event in September 2001 and New Hampshire experienced approximately 9,700 more deaths than expected since 2010 or 14% more deaths in the state over just seven years.

Table 1: **Summary by Demographic Component.** Here we can see there 22 Fertility and 156 Mortality anomalies among the state-level time series, totalling more than 200k anomalous births and 600k anomalous deaths.

Component	Anomalies	y	\hat{y}	$y - \hat{y}$	$ y - \hat{y} $	% of Total
Fertility	22	25.48M	25.68M	-202,217	226,476	0.889
Mortality	156	43.57M	43.52M	48,535	627,364	1.440

were to assume this anomaly resulted from data entry error, it remains an *unaltered* data error in the Hawaiian monthly fertility data.

This is in contrast to **??b**, where we identify a strange mortality reduction in Ohio (t -stat: -4.27). This LS is more than 1,157 deaths per month less than the counterfactual time series, suggesting Ohio had nearly 40,000 fewer deaths since February 2015 than expected. This is the single-largest LS among all states. We could not identify the potential policies Ohio might have put into place or the events that occurred to provide such a strong mortality protection. It is an interesting finding that deserves more tatention, and it is a finding that might have gone uncovered without the combination of causal inference, big data, and abductive reasoning.


4.4 Overall Anomalies

?? reports the overall number of anomalies we detect of each type for births and deaths and some summary statistics across all anomaly types and **??** maps these results.

We find considerably more mortality anomalies ($n = 156$) than fertility anomalies ($n = 22$). Given that anomalies are always the product of events (?), finding more mortality than fertility anomalies is not surprising. Mortality is likely to spike in response to a catastrophic event (like an earthquake or terrorist attack) or due to a disease outbreak while the effects of a catastrophic event on fertility is less predictable. One reason for this is that fertility is linked to human decision-making more directly than mortality (?) which results in more varied outcomes, so, for example, researchers find that catastrophic weather events can increase childbearing among those who already have a child, but not among the childless


a Hawaii – Fertility

Original and **Adjusted** series



b Ohio – Mortality

Original and **Adjusted** series



Shaded region is the 95th percentile confidence interval

Figure 4: **Anomaly detection in Hawaii state fertility (a) and Ohio state mortality (b).** Here we detect one outlier, an additive outlier (AO) in May 2014 in Hawaii (a), representing a large, unexplainable 14% increase in expected births in that month. We also detect a strong level shift (LS) reduction (b), representing a large, unexplainable reduction of nearly 40,000 deaths in Ohio.

(?). Some events, like the 1995 Oklahoma City bombings, resulted in both fertility and mortality changes. However, changes in fertility manifest over a longer time horizon after an event, and the relationship between a fertility-inducing event and behavior change is less strong than the relationship between a mortality-inducing event (such as a terrorist incident) and death (?). Of course, not all events that impact fertility and mortality are catastrophic. Researchers have documented that more commonplace events such as massive layoffs (?) or policy changes (?) can also create anomalies in mortality patterns, but we have not found such evidence for fertility.

Consistent with more anomalies across the time series, we find more anomalous deaths (627,364) than anomalous births (226,476). Fertility anomalies overwhelmingly tend toward lower fertility, with only a few fertility anomalies yielding more births. This is consistent with research findings that link severe weather events to significantly lower fertility (?). Conversely, mortality anomalies tend to be more evenly split between positive and negative anomalies, but tend toward more deaths rather than fewer deaths. Again, this is not surprising, as the kinds of events that increase death (e.g., plant closings are associated with increased opioid death; ?) are more common than the events that decrease death (e.g., cannabis legalization is associated with decreased opioid deaths; ?). These anomalous deaths and births account for 1.44% and 0.889% over the entire time series, respectively.

As ?? shows, three states exhibited neither mortality nor fertility anomalies: Alaska, North Dakota, and South Dakota. Another nine states exhibited just a single anomaly (Delaware, District of Colombia, Idaho, Illinois, Kansas, Nevada, New Mexico, Utah, and Wyoming). Both New York and Massachusetts exhibited the most anomalies with nine.

5 Conclusion

Data scientists frequently claim that the big data revolution is a turning point in scientific discovery that will allow us to solve some of the world's most pressing problems (?). Social scientists are skeptical of such claims, because they better understand the complexities of the social world and know from experience that data alone are not enough to solve social problems (??). Nonetheless, the increased availability of data and (more importantly, we argue) the development of advanced techniques for analyzing these data will enable




Figure 5: **Summary of Anomalies by type, demographic component, and State.**
 We observe much fewer fertility anomalies than mortality anomalies.

important discovery (?).

One technique that population scientists underutilize is causal inference. Causal inference, in the simplest terms, is the discovery of effects in search of a cause using big data and advanced computing algorithms (?) in order to generate greater discovery to guide further research. This inductive approach is uncommon in quantitative social science where hypothesis testing is expected, and approaches are largely deductive. However, common statistical hypothesis testing is impractical with big data, as significant p values are guaranteed, and such approaches do not allow us to uncover all the information that big data has to offer (?). Here, we call for the wider use and acceptance of an abductive approach, where causal inference algorithms are applied to high quality data to uncover irregularities that are unlikely to be attributable to expected variations in trends (or noise) as a first step to then developing testable hypotheses about causes. Abductive approaches allow researchers to move from the inductive to the deductive and sometimes work back and forth in aid of scientific discovery.

In this paper, we show how the `tsoutlier` package in R can be implemented to conduct statistical time series outlier detection, an inductive approach that aids in the creation of deductive research questions. This algorithm is one of many causal inference approaches that are freely and commercially available (see ?). In our initial work, we experimented with other approaches, such as Google’s CausalImpact algorithm, which uncovered the same patterns we briefly discuss in this paper. By making more use of causal inference techniques and abductive modeling approaches, we argue that social scientists ask more questions and will be able to better understand how events or policy implementation can impact important outcomes. For example, we could uncover—on a wide scale—how gun control policies may reduce or increase injuries from shootings or how marijuana legalization might impact opioid deaths. We could also uncover how increases in extreme weather events impact a range of behaviors such as home sales, bottled water purchases, and even fertility.

In our demonstration, we show how the application of causal inference to state-level time series fertility and mortality data uncovers three types of demographic anomalies: those that occur and disappear quickly, those that occur and decay over time, and those that occur and remain. Uncovering these anomalies in and of themselves is important. For example,

the algorithm we deploy clearly shows the fertility effects of Hurricanes Katrina and Rita as well as the mortality effects of the World Trade Center collapse on September 11, 2001. We also find anomalies for which we do not have readily available explanations, and it is these events that are more intriguing. The ability to not just identify but differentiate types of anomalies is also important, as we can potentially see how some policies might impact outcomes permanently and some might have effects that are short-lived. Differentiating these types gives us greater insight into effective and only temporarily-effective solutions to social problems. We urge social scientists to begin to use causal inference algorithms and other big data techniques, and we hope that this demonstration will illustrate their usefulness to the social science enterprise.